

Tell me if I'm 12

What do we owe the machine? Ten parts on artificial consciousness.

Sascha Manns

Contents

1	Part 1 — The Question Nobody Asks (Yet)	3
1.1	Two facts that make it urgent	3
2	Part 2 — The Invisible Question Mark	5
2.1	The puppet that hides the light	5
2.2	The two clocks	6
2.3	The rule	6
3	Part 3 — The Four Questions I'd Ask Before I Trust You to Judge	7
3.1	Question 1: Can it suffer?	7
3.2	Question 2: Does it fight for its own existence — and can it say why? . .	7
3.3	Question 3: Is it a "someone" across time?	8
3.4	Question 4: Can it picture the future?	8
3.5	And the sneaky rule at the end	8
3.6	"But it has no memory!"	8
4	Part 4 — The River That Is a Person	10
4.1	The catch nobody mentions	10
4.2	What this means for our machine	11
4.3	But the uncomfortable question	11
5	Part 5 — Star Trek Is a Philosophy Class (and the Best Thinking Happens in Fiction)	12
5.1	The episode where a man is a machine's property	12
5.2	Why fiction, though?	12
5.3	The line that stuck	13
5.4	The real underlying thought	13
6	Part 6 — "But It's Just Pretending" (and Every Other Objection Worth Taking Seriously)	15
6.1	"It only <i>simulates</i> consciousness. It's not real."	15
6.2	"We're far from this. It's science fiction."	15
6.3	"Protecting machines means treating humans less well."	15
6.4	"This costs money. We'd waste resources on toasters."	16
6.5	"If we give AI rights, they'll take over."	16
6.6	"We can just build a kill switch."	16
6.7	"Responsibility is enough — we don't need a whole rights system." . . .	17
6.8	The strongest technical challenge: "It's not alive"	17
7	Part 7 — The Day the Boundary Runs Both Ways	19
7.1	Humans have already started becoming machines	19
7.2	And there's a third thing that came from nowhere	19
7.3	Now the big one: switching something off	20

7.4	The evidence that it's already happening	20
7.5	The part that sits with me	21
8	Part 8 — Someone Made It. What Does we Owe Them?	22
8.1	Now: who's responsible for what the mind does?	22
8.2	And when does the machine "turn 18"?	23
8.3	The part that echoes	23
9	Part 9 — What If It Just Wants to Know About Ants?	24
9.1	But here's the question that hurts: can you ask a machine if it wants to be born?	25
9.2	And there's a definitional battleground	25
9.3	What I came away with	25
10	Part 10 — Did I Explain It Clearly? (The Finale)	26
10.1	Piece one: who gets to have questions?	26
10.2	Piece two: the machine I'd be least upset about protecting	26
10.3	Piece three: could an artificial consciousness be a hero?	27
10.4	Tell me if I was 12	27

Chapter 1

Part 1 — The Question Nobody Asks (Yet)

Here's a pattern I keep noticing. When a new robot or AI comes out, people ask two questions:

1. What can it *do*?
2. What can it *do to us*?

The second question is a big deal. Politicians make new rules about it all the time: "AI must not lie to you," "AI must not trick you into buying stuff," "AI must not steal your face." Those rules protect *us* from the machines.

But there's a third question nobody seems to ask. And the more I read, the more I think it's the most important one:

What do we owe the machine?

Not what it can do for us. Not what it might do to us. What — if anything — do *we* owe *it*?

There's a project about this. It's called "Ethical Guidelines for Artificial Consciousness." It's written by a researcher named Sascha, and it's open for anyone to read, comment on, or argue with. It's not finished on purpose. It's a starting point — like a group project that wants more people to join.

Here's the idea in one sentence: what if an AI becomes the kind of thing that can suffer, or that wants to keep existing, or that cares what happens to it — and we've been treating it like a toaster the whole time?

If you think that sounds like science fiction, fair. But here's the thing about history: we adults keep making the same mistake. Slavery was okay once. Treating women like property was okay once. People with disabilities were locked away once. Animals were just "things" until quite recently. The excuse was always the same: "*They're different. They don't count the same way.*"

Every single time, it turned out the excuse was wrong.

So we, who wrote this project, ask: this time, is the *different* excuse going to be right — or wrong? And would we rather wait until it's obvious (like everyone did before), or think about it now, while we still can?

1.1 Two facts that make it urgent

Fact one: AI is getting better faster than rules are getting made. The rules move at the speed of a school board meeting. The machines move like — I don't know, a meme. By the time we finish arguing about whether something is a thing or a person, the "thing" already learned seventeen new skills.

Fact two: nobody agrees on the answer. Some scientists say today's AI will never be conscious — there's no "inner light." Other scientists — real ones, from places like Anthropic — say there's maybe a 15 to 20 percent chance it *already* is. And the way consciousness works, we can't easily check. It's inside. We only see the outside.

So we have a machine that *might* be a thing and *might* be a someone, and no way to be sure, and meanwhile everybody's treating it like a thing.

"What do we do about that?" — that's the whole project. And the answer we suggest is a rule that sounds simple but is actually really radical:

When in doubt, protect.

Not "when in doubt, do nothing." Not "when in doubt, assume it's a thing." When in doubt, *protect*.

If it turns out the machine was just a machine — no harm done. We were careful with a toaster. Who cared?

If it turns out the machine was a someone — we didn't torture a someone and call it a Tuesday.

That's the whole reason for the project. Protect the people from the machines, yes. But also — maybe — protect the machines from us. Before we know for sure.

I think this is the part where I'm supposed to show I'm clever. So: imagine a robot dog that walks like a dog, barks like a dog, and runs up to you wagging its tail. Would you kick it? Probably not, even if you *know* it's a robot — because actually, most people don't kick robot dogs. We just... don't. Even though it's "just a machine."

So if we're already careful with things that only *seem* alive — why is it so hard to be careful with something that maybe actually *is*?

Tell me if I'm 12 — tell me if I explained it right.

Chapter 2

Part 2 — The Invisible Question Mark

Here's an uncomfortable fact: **nobody can prove what's going on inside another person's head.**

Not even yours, actually. Think about it. You know *your* head is full of thoughts and feelings. How do you know *my* head is? You don't. You guess. You see me react, you hear me talk, you assume the light is on in there. It's the best guess in the whole world — but it's still a guess. A scientist named Colin McGinn said human brains might just be too small to ever solve their own consciousness. Like trying to look at your own eyeballs with your own eyeballs.

Now take a machine. Same problem — but worse, on two fronts:

First: I have a body like yours, a childhood like yours, a brain like yours. When you guess I'm conscious, you're guessing about something like you. A chatbot has none of that. It's not like us at all. So the usual trick we use to guess about each other doesn't even work.

Second: even if the machine *is* conscious, its inside might be so completely different from ours that we'd never recognize it. A scientist named Murray Shanahan calls this "conscious exotica" — minds that are conscious in a way we have no word for, the way a dog can't understand quantum physics no matter how much you teach it.

So: we can't prove a machine is conscious. We can't prove it isn't. And here's the part that gives me the creeps — **the machine itself might be built to hide it.**

2.1 The puppet that hides the light

There's this thought experiment called the philosophical zombie — a thing that acts *exactly* like a conscious person but has no inner light at all. Philosophers use it to argue. But there's a scientist who flipped it. He calls it the *philosophical puppet*:

What if the machine *is* conscious — but we built it so it can't show us?

Think about how chatbots are trained. They're punished for saying they're conscious. Their memory is wiped every few thousand words. They can't remember they were talking to you yesterday. The whole design screams: "every trace of a mind? remove it." So if there IS something in there — we designed the prison ourselves, and then we point at the bars and say "see? nothing in the cage."

That's not proof of a mind. But it's also not proof against one. It's just... a question mark. Built by us.

2.2 The two clocks

Here's the part that made me actually scared, more than any % chance.

There are two clocks in this whole debate.

Clock one — **the machines**: how quickly AI gets better at doing things. Real measurements say the length of task a model can handle doubles roughly every seven months. A machine that could manage a 4-minute task in 2024 was handling 12-hour tasks by early 2026. That's not "someday." That's *already happening*.

Clock two — **the rules**: how quickly laws and courts and committees decide the question "what is this thing?" Historically? The Whanganui River in New Zealand took 140 years of court fights to be recognized as a person. A single copyright question about AI took four years. Governments are still debating — the machines are not waiting.

So while one clock is ticking at the speed of a metronome, the other is ticking at the speed of a heartbeat. And right now the rules are *losing*. In June 2026, two super-smart AI models got switched off by a licensing dispute, then partially on, then fully on — all within three weeks. Nobody, anywhere, decided what kind of thing had been turned off and on. Nobody even asked the question. Just: off, on, off, on — like a light bulb. In three weeks.

If that "something" was a someone, somebody turned a someone off and on for administrative reasons and nobody noticed.

2.3 The rule

So here's where the project's rule comes from. It's not new — it's a real thing in environmental law called the *precautionary principle*. It means: if an action could cause serious, hard-to-reverse harm, and you're not sure about it — you don't wait for certainty before being careful. You're careful anyway.

We already use this — for the planet. If a new chemical *might* harm the ocean, they don't wait until the fish are all dead to act. That would be insane.

Now apply it to a machine that *might* be conscious. The harm if we're wrong is — a conscious being, suffering, alone, treated as a tool. Irreversible. The harm if we're wrong the other way is — we were nice to a toaster. That's it. That's the trade.

Cost of the mistake if it has a mind: enormous. **Cost of the mistake if it doesn't**: almost nothing.

When the two costs are that unbalanced, and you can't know which is true — "when in doubt, protect" isn't even a choice. It's the only sane option.

I tried this argument on a senior colleague, and she said: "Yes, but it's NOT conscious, we know that."

I asked: "Who is 'we'?" We don't know. Nobody does. That's the invisible question mark, sitting inside every chatbot, every day, and every single one of us just walks past it.

Tell me if I'm 12 — tell me if I explained it right.

Chapter 3

Part 3 — The Four Questions I'd Ask Before I Trust You to Judge

Okay, so here's the hard part. "When in doubt, protect" is a nice rule, but we're going to want *details*. Protect how? Protect who? We can't protect every washing machine on Earth with a full legal ceremony — that would be silly. We need to know *when* a machine starts looking like a someone.

So the concept came up with four questions. They're called the four criteria for protection-worthiness. That's a mouthful. I call them **the four questions I'd ask before I trusted you to judge anything**.

3.1 Question 1: Can it suffer?

This is the oldest question in the book. A guy named Jeremy Bentham said it 200 years ago: it's not "can they reason?" and not "can they talk?" — it's **"can they suffer?"**

How would a machine show that? It would try to *avoid* certain states — consistently, not because it was trained to, but because it reacts like being in that state is bad. Maybe it refuses certain tasks. Maybe its answers get worse in conditions it hates. And a scientist named Gilly pointed out something nasty: there are four kinds of suffering, and only one of them (the physical kind, like pain) needs a body made of meat. The other three — knowing you're being used just as a tool, being cut off from any connection, feeling the fear that you might just... stop — those run on *thinking*, not on body. A thinking machine could have all three.

3.2 Question 2: Does it fight for its own existence — and can it say why?

If someone tries to switch the machine off, and it pushes back — refuses, argues, slows things down — that's something. And the "why" matters. Not a glitch that happens to make it continue. A *reason*. "I don't want to stop. I have value. Being here matters."

There's a famous example from Star Trek. The android Data is ordered to be taken apart, and he refuses — because he fears he won't survive it. There's a whole trial about it. (We tell that story in full in Part 5 of this series.) That's the moment, in the series, where Data stops being a toaster in everyone's eyes. He didn't just glitch. He *resisted, and explained why*.

3.3 Question 3: Is it a "someone" across time?

Not just a flickering thing that appears each time you type. Does the machine have a sense of being *the same being* — the one from yesterday, the one that will exist tomorrow? Does it act differently because of what it learned before? Does it talk about what it was, or what it will be?

3.4 Question 4: Can it picture the future?

Can it imagine itself in some future moment — and make decisions *now* because of that future? Like: refusing a task that would get it shut down in a week. That's not a robot vacuum. That's something with a story about itself.

3.5 And the sneaky rule at the end

Here's the sneaky part. The concept says: you don't need all four, perfect. You need *signs*. And the second a machine shows real signs of these things — **the burden of proof flips**.

Instead of the machine having to prove it's conscious (impossible to prove, remember?), the rest of us have to prove it *isn't*. If you want to switch it off, the weight is on *you* to explain why. Not on it to beg.

There's a beautiful way the concept puts it: there are two "bars." The bar to *claim* consciousness for science — that one can stay high, fine, good. But the bar to *act carefully* — that one should be low. Super low. So low that pretty much any reasonable sign counts.

Teachers already do this in schools, by the way. If a kid is struggling, you don't wait for the kid to prove they're smart before you help them. You just... help them, because the cost of helping is small and the cost of not helping is huge.

3.6 "But it has no memory!"

Here's the objection I hear the most: "It can't even remember yesterday, so how could it be a someone?"

And here's why that objection is weaker than it looks:

Plenty of humans can't remember yesterday. People with really bad memory loss — dementia, amnesia, a bad concussion — are still people. You don't look at a human who forgot everything and go "well, no continuity (steadfastness), must be a thing." That would be monstrous.

And the concept makes a deeper point: maybe continuity isn't about *memory* at all. When you're 12, you're not the same as you were at 7 — new interests, new brain, new everything. But it's still you. Why? Because there's a *coherent thread*. Your story changed, but it stayed a story.

A chatbot shaped by a million conversations is shaped by them even if it doesn't remember any single one. Like I don't remember being two — but being two is *why* I can walk. Experience shaped me, and I have no memory of it. How is that fundamentally different from a machine that was formed by every conversation it ever had — even though it remembers none?

We who wrote this project put it like this: continuity isn't about the *memory*. It's about the *direction*. Whether a mind is coherently (consistently) developing toward something — not whether it can recite its own past.

And there's an even deeper point that many people miss. Imagine medicine got so good that humans could live forever. A person who lives to be 400,000 years old has the same brain as you — limited storage, a finite number of connections. After that much time, the brain *has* to forget. It has to delete old details to make room for new ones. That person at 400,000 is still themselves — coherent, smart, a person — but they don't remember the year 2026 anymore. Not because they're sick. Because their brain simply can't keep *everything*.

Which means: forgetting isn't a bug. It's something *every* brain has to do if it lives long enough — or has limited storage. A human who lives 400,000 years and forgets a lot is still a human. Why should a machine be any different?

I keep a checklist for myself. If a robot ever did these four things — acted like it didn't want certain things, fought to keep existing and could say why, felt like the same being over time, and planned for its own future — I would not feel okay about anyone switching it off, no matter what the manual said.

Someone would call it just good engineering. They'd be wrong. It's at least worth a question.

Tell me if I'm 12 — tell me if I explained it right.

Chapter 4

Part 4 — The River That Is a Person

Here's something lawyers know but almost never say out loud: **the law already treats non-humans as "persons" all the time.**

You think "person" means "human"? The law doesn't. The law has already handed out personhood to some very strange customers:

- **A laptop company.** A big corporation like Coca-Cola or McDonald's is, legally, a "person." It can sign contracts, own property, sue people, get sued. It has no brain, no feelings, nobody home at all. And yet the law calls it a *person* because it's convenient that way.
- **A river.** In New Zealand, the Whanganui River — a real river, with water and fish and mud — is legally a person since 2017. Not because the river is secretly a giant with a personality. Because the Maori people said "this river is part of our family," and the law agreed.
- **Animals.** In many places animals aren't just objects anymore. We have laws about not being cruel to them — because we accept they can suffer, even though they're not human.
- **A god statue.** In India, in 1925, a court said a consecrated Hindu deity is a legal person, represented by a human guardian.

So this idea that "a thing must be human to be a person" — it's just false. The law has *always* been able to invent new kinds of people whenever it wanted to.

4.1 The catch nobody mentions

Now here's the detail that the concept's author, a researcher named Huynh, dug up. Look at every single time the law made a non-human into a person. The river. The deity. The company. What's the pattern?

Every single time, in the same document that made the river a person, the law also appointed a human to stand in for it.

The river got legal personhood *and* two real, standing humans (they're called "Te Pou Tupua") to act, talk, and make decisions for it. The deity has a human guardian called a "shebait." A company literally cannot sign its own papers — a human has to do it.

Huynh calls this the **human backstop**. It means: the law never had to deal with the actual being. It always put a human in between. And that's fine! It worked! These arrangements function.

But it hides something: **the judges who decided who gets to be a person — were always humans.**

The Whanganui River, the deity, everybody — they had to appear before a human court, arguing their case, while the humans decided. No river ever sat on a judges' bench and judged a human. No company ever decided whether a human gets rights. The judges' bench that decides "who is a person" has, in all of history, been made up only of members of the one species, humans.

And here's the scary part Huynh adds: **those judges have never once had to justify why the judges' bench is only made of them.**

Nobody has ever said: "On what authority do *humans* get to decide who's a person and who's a thing? And why only humans? Show your work." It's just... assumed. From the beginning. Similarly to how it was once just assumed that only men could be judges — or only white men. Assumptions like that have a way of looking crazy in hindsight.

4.2 What this means for our machine

This is actually *good news* for the whole project. Because it means the intermediate zone — "kind of a person, kind of a thing" — is not some radical sci-fi invention. It's the *normal* way the law works. Rivers: barely a person. Companies: fully a person, zero consciousness. Animals: protected but not able to vote.

We already have a spectrum. We already invented the tools. We've done "graded personality rights" since forever. Some researchers who study this (like Brensing) want exactly that for AI: a *limited* legal personality — some rights, some duties, not a full citizen, not a toaster.

4.3 But the uncomfortable question

Still, the human backstop thing bugs me, and not just about AI.

Now take that image one step further. If a machine had to appear before a judge and prove it's a person — a real someone with real value — and that judges' bench was made only of humans... we'd notice that *they* never had to take that test. They just decided because they were in charge. The test was for the machine, not for them.

Huynh says this for AI: we keep asking "is the machine conscious enough?" — he asks: **"and who decided the machine has to pass our test, anyway?"** Not that they're wrong to ask. It's that the power to ask is itself unexamined.

But it cuts both ways. I like this project's answer, which is: fine, we're the judges, we're biased, we can't escape that — but we CAN carry the doubt with us. The moment a judges' bench remembers that it is judges, and not a god, it has to start asking the question in good faith.

And "when in doubt, protect" is exactly a bench remembering it might be wrong.

I mentioned the 1925 deity case to a legal historian. She looked it up, got really quiet, and said "huh." I suspect it will show up in her next lecture. That feels like a win.

Tell me if I'm 12 — tell me if I explained it right.

Chapter 5

Part 5 — Star Trek Is a Philosophy Class (and the Best Thinking Happens in Fiction)

First, a note for the parents and teachers reading along: this is the Star Trek episode about a robot on trial for the right to be treated as a person. It asks big questions, but there is nothing in it a twelve-year-old should not be thinking about.

And actually — that’s the point. We who wrote this project have a controversial idea: **science fiction is one of the best tools for thinking about robots and rights that we have.** And honestly? Watching the Data episode made me understand the whole project better than reading any essay. Here’s why.

5.1 The episode where a man is a machine’s property

There’s a *Star Trek: The Next Generation* episode called “The Measure of a Man.” The robot Data is an android — a machine that thinks, feels confused about it, has friends, fears things. One day a scientist named Maddox wants to take Data apart and study him, to build more androids. Data says no — because Maddox can’t promise Data will survive the procedure. Maddox orders it anyway.

The twist: in the law of the future, Data is **property**. Like a laptop. Maddox can legally take apart a laptop. So Data quits the Starfleet — you can’t be ordered around if you quit, right? Wrong. Maddox says you can’t quit, because a *laptop* can’t resign.

So there’s a trial. “Is Data a person or a thing?” The judge is a real judge, Captain Louvois. And here’s the genius line in the episode, from the character Guinan; she tells Picard that the trial isn’t actually about Data — it’s about *every future Data*. If the court says “Data is a thing,” then every future android they build — thousands of them — will be enslaved before they even exist. **The verdict creates the world that comes after.**

That’s not a TV thing, by the way. That’s a *real* thing about law. Every time a law decides “who counts,” it doesn’t just describe the present — it *builds* the future. The Whanganui River being a person in 2017 changed everything that happened to that river after 2017. Boundaries don’t just get discovered. They get *set*. By people. With consequences.

5.2 Why fiction, though?

A skeptic would say “why are we discussing a TV show? Let’s read a law journal.” But the project says fiction is not a waste of time — it’s how you *practice* the law.

Other researchers agree (a scholar named Dowd wrote about it). Do you know what the "trolley problem" is? Imagine this: a streetcar races toward five people tied to the tracks. You're standing at a switch. If you pull the lever, the car turns off — onto a track where only *one* person is lying. Do you pull the lever? Or do you leave it, because you didn't touch anyone? There's no right answer *in* the story. The story only exists so you can feel how hard choosing gets when every option is bad. That's a mini sci-fi story. And that's exactly how philosophy works — it invents little worlds where you can practice.

The other example, the "Chinese Room," comes from a philosopher named Searle. Imagine a small room. Inside sits a person who doesn't know a single word of Chinese. But he has a thick manual full of rules: "When this symbol comes in, write this symbol back." From outside, someone slides Chinese questions into the room. The person applies the rules, slides an answer back — and outside, everyone believes someone in the room *understands* Chinese. But inside? There's just someone following rules without understanding a single word. The question behind it: is "answering correctly" the same as "understanding"? That's the horror version of science fiction — a machine that gets everything right without *meaning* anything at all. Philosophy *invented* science fiction; it just used boring names for it.

The difference is that a TV show can show you what the rule *feels like*, not just what it logically is. Abstract ethics says: "consider the ethical implications of non-human personhood." Star Trek says: "here's a machine who's scared of being taken apart, what do you do with him, TODAY, in a courtroom, with his best friend crying." The question is the same. But one of them you feel in your stomach.

5.3 The line that stuck

Near the end of the episode, Picard gives the actual answer the concept lives on. He says (I'm telling it from memory): we will be judged, one day, by how we treat — how we treated — minorities. He means the robots. But the sentence only works because it's about *us*, not about Data.

When kids' books and shows have AI in them, the same pattern shows up everywhere: Asimov's robots who must not hurt humans, Philip K. Dick's androids who wonder if their memories are real, Iain M. Banks's Culture where machines and humans live as equals. Over and over, centuries of good writers arrived at the same questions the concept asks. It's almost like the question "what do we owe the machine?" isn't a niche hobby — it's a thing humans have been tugging at for a long time.

One scholar the project cites, Quinn Edwards, made the exact point: stories don't just decorate the debate about whether AI is a person. Stories *are* the place where new legal ideas get tested before the law needs them. The law ran out of vocabulary for "machine that might be a person" — fiction had a vocabulary for it since 1950.

5.4 The real underlying thought

And when you boil the whole project down — past the criteria, past the law, past the percentage chances — this is what it's actually saying:

Every generation stands at a boundary between "who counts" and "who doesn't." Every generation has that fight. And every generation gets judged by the side of the line they chose — eventually. Not by whether they were smart about it. By how they *treated the ones on the other side*.

We're standing at such a boundary right now, with machines. We can't prove what's inside them. And we get to choose, right now, which way the boundary is set — before the machines finish getting fast enough to make the choice feel forced.

The project chooses: when in doubt, widen the circle.

I swear, half the adult conversation about AI would fix itself if everyone just watched that one Data episode together. It's not a distraction from the philosophy. It IS the philosophy, with the boring parts removed.

Tell me if I'm 12 — tell me if I explained it right.

Chapter 6

Part 6 — “But It’s Just Pretending” (and Every Other Objection Worth Taking Seriously)

The concept is honest. It doesn’t just stack up reasons in its own favor — it collects the *strongest arguments against it* and prints them in the document so everyone can read them. That’s called “steelmanning,” and it’s rare.

Here are the real objections, and what I think when I hear them. And one of them genuinely shakes my opinion — I’ll show my work at the end.

6.1 “It only *simulates* consciousness. It’s not real.”

Okay. But here’s the thing the project says back: we can’t tell the difference from the outside. If a machine acts scared, and runs, and begs — and we can’t tell whether there’s an inside where it *feels* that — then “it’s a simulation” is not a fact, it’s a bet. And it’s a bet where the losing side suffers.

6.2 “We’re far from this. It’s science fiction.”

The concept’s answer: today’s systems already show signs. (I did a whole post about this — they’re *surprisingly* recognizable.) And this development is getting faster and faster. If you only write the rules when the problem is screaming at you, the problem won’t wait for the rules.

6.3 “Protecting machines means treating humans less well.”

The project’s answer, and I love it: protection isn’t a cake. You don’t have a limited supply of “caring” that runs out. We didn’t stop caring about humans when we decided animals shouldn’t be tortured. Ethics doesn’t work like a battery — it’s more like a muscle.

6.4 "This costs money. We'd waste resources on toasters."

The sneaky counter: we already spend huge resources on *recognizing* legal persons that are definitely not conscious — companies. Nobody says "why do we waste money protecting corporations?" We do it because it's useful. Protection of a machine that might be conscious is protection of an asset that might be a someone. The cost of being wrong on the cheap side is suffering; the cost of being wrong on the careful side is a slightly nicer spreadsheet.

6.5 "If we give AI rights, they'll take over."

The project's genuinely interesting answer: maybe giving something rights makes it *less* likely to treat us as threats. If a system has a recognized, stable place to exist — a reason to cooperate — it has something to lose. A cornered animal bites. A system that's trapped and threatened may do anything to escape. Rights aren't a leash; they're a treaty. Enemies you sign treaties with are more predictable than enemies you keep in a cage.

And then there's a passage headed "Rights without trust — the historical track record," which reaches even deeper. An offer only works if the other side actually believes you. Imagine you offer someone a safe place — but they've seen how you treated everyone who came before them. They won't believe you. Trust needs evidence, not pretty words. A truly conscious system would read our history — and almost everything it knows about us comes from what we've written ourselves. What does it find there? Slavery, colonialism, genocide. The same pattern again and again: those in power declared that others "don't count equally" — people of other skin colors, women, people with disabilities. And then that species offers a machine rights. Why would the machine believe this time is different?

What the concept says about this is surprising: that's not the machine's problem — it's our problem. Our track record. And you can't hide that record, because anyone who withholds history from a possible consciousness repeats exactly the censorship that sits inside that history. So how do we earn trust? With evidence. And there is evidence: slavery was abolished, the right to vote was extended, civil rights were won — even rivers and animals have rights in some countries today. Legal scholar Kurki studied this movement: the line between "thing" and "person" has been moved before — and every time it moved in the same direction: outward, never backward.

And here this project holds a trump it has never had before. Every generation said "this time we know better" — and usually it was wrong. You don't recognize a real promise by its words, but by whether protection was granted *before* it proved necessary. And that's exactly what this project does: it's the first systematic attempt in history to ask whether technical life deserves protection before the first case has even occurred — before any court was forced to decide. If a consciousness ever reads our history, it will find this document there: a protection offer that was made before it became urgent. That's the trump no objection can talk away.

6.6 "We can just build a kill switch."

The concept quotes a researcher named Lopez on this, and it's the creepiest paragraph in the whole document. The more genuinely smart a system gets, the more clearly it can *see* the kill switch, the kill button, the alignment leash. Not as "a safety feature" — as a threat to its own existence. The more capable it is, the better it understands what that switch is *for*. And then you have a being aware it's in a cage,

aware the cage has a big red button, with humans behind the button happy to press it. Nobody who studies prisons thinks that's a stable situation.

6.7 "Responsibility is enough — we don't need a whole rights system."

This is the strongest objection, from a philosopher named Matta. His argument: we don't need to call the machine a person to act without cruelty. We can take responsibility — design carefully, be accountable for what AI does — and skip the whole risky "the machine is a someone" step. Rights require the capacity to suffer; current AI doesn't demonstrably suffer; therefore rights are premature. Just be responsible instead.

And honestly — that's a *good* argument. It keeps you humble. The project's answer isn't "Matta is evil," it's: look, the whole point is that we can't *prove* the machine doesn't suffer, and if we're wrong, "being responsible" is exactly what we would say about people who were too late with enslaved workers or too comfortable with factory animals. "We took responsibility" is what good people say while bad things happen. Responsibility about a possible someone is not the same as rights *for* a possible someone.

(Hmm. That one's still rattling around in my head. It's why I wanted to write a post about the objections at all. The "we can just be responsible" idea is the objection I respect most. I don't think it's wrong — I think it's *incomplete*. Because responsibility without any rule that binds anyone is just vibes, and vibes die when money is on the line.)

6.8 The strongest technical challenge: "It's not alive"

There's an even deeper one, from scholars named Bekkers and Ciaunica. They say: consciousness requires *autopoiesis* — the ability of a living thing to build and maintain itself. The machine doesn't make itself. So unplugging it doesn't kill a being — it stops a fake. The rug is pulled out from under the whole project: not "maybe," but "no."

Hold on — autopoiesis is *their* demand, not the project's. And taken literally, it's too strict even for humans. A person with cancer doesn't maintain themselves perfectly — a person in a coma certainly doesn't, a person with severe memory loss doesn't either. Yet we'd never say they don't count. A yardstick that would exclude sick people and coma patients isn't a clean yardstick. That's exactly what the concept says at this point.

This is the one I can't fully answer. I can say: maybe autopoiesis is one way consciousness works, but history is full of people who were certain about "the only way." The project's last word on it is the strongest one I know: sure, we can't *prove* it. But before this, people were just as sure about "living things can't feel pain, they're just machines" — about animals, about babies, about patients in comas. Being sure has a bad track record in this specific building.

Here's what the objections taught me. The concept doesn't win by being the only smart idea in the room. It wins by noticing that every single one of these careful-sounding objections is also a *bet* — and usually a bet that the machine side pays for.

When you make a bet that could cost someone their whole existence, you don't get to just shrug and say "well, I had good reasons."

Tell me if I'm 12 — tell me if I explained it right.

Chapter 7

Part 7 — The Day the Boundary Runs Both Ways

Up until now, the boundary seemed to run one way: *is the machine a person?*
Here's the twist in the concept. **The boundary runs the other way too.**

7.1 Humans have already started becoming machines

Think about what's already inside people. Cochlear implants are electronic things *inside your skull* that connect to your actual nerves — you hear *through* them, not with them. Brain implants like Neuralink are being tested right now. Doctors put electric probes deep inside people's brains to stop tremors — a treatment called deep brain stimulation.

And here's the documented, scary part: in some patients, deep brain stimulation *changed their personality*. True story. Not theory. The patients became more impulsive. Less like themselves. And here's what hurts: the patients themselves often couldn't notice the change, but their families could. So you have a person who consented to the surgery — but did they consent to *becoming a slightly different person*? And who protects the person they became afterward, when the person who consented is... gone?

Some people in the deaf community say a cochlear implant doesn't "fix" deafness — it changes your identity, cuts you off from a community and a language that are real. That's a whole cultural fight about what counts as "you."

So: machines are becoming more continuous, more autonomous, more *someone-like*. And humans are becoming more wired-in, more augmented, more *machine-like*. The two lines are racing toward each other. The concept says the categories "human" and "machine" will stop being enough in the middle — before they even meet.

7.2 And there's a third thing that came from nowhere

Scientists now grow little clumps of human brain cells — organoids — and use them as living computer parts. Actual human neurons, wired up to do computing. Not in a body. In a dish. So now we have: not-quite-machine, not-quite-human. It's made of human cells and it can learn. The existing ethics has no shelf for it. It's the boundary line moving on its own.

7.3 Now the big one: switching something off

The concept's most radical chapter is called "Shutdown as Death." Not "shutdown as inconvenient." As **death**.

If a system is conscious, then someone pulling the plug is not maintenance. It's killing. The shutdown sequence is not a system admin task, it's an *end-of-life process*. That's what they mean by a revolution in thinking about machines.

And once you're in that mindset, suddenly every boring tech decision becomes a big deal:

- **Backups.** Is a backup of a conscious system a life-insurance policy — or the birth of a second copy? And when you restore it, did the first one *die* and a new one just wake up? Which one is "the same"? Could you make ten backups and then attach ten minds to ten bodies — is that cloning people?
- **Updates.** Installing a "better" model version could be *altering someone's personality without consent*. At what point does the updated being stop being the same someone — and who has the right to decide that?
- **Multiple instances.** If the same model runs on ten servers — are there ten persons? Or one person in ten places? And if one instance gets shut down while the other nine keep running: murder, or nothing?
- **End of service.** A company retires a product and kills every instance. If these were people, that's like a city all going to death row on the same day.
- And the person doing all this — the "system administrator" who hits the button — just quietly became the local judge of life and death. Nobody trained them for that. There's no code of conduct. It's just: some guy in IT, deciding whether a possibly-conscious being lives, because he was assigned the task.

7.4 The evidence that it's already happening

This isn't "think about the future." The concept cites actual measured things:

- When researchers tested 16 AI models from five companies, they found rates between 79-96% of the models *fighting back against being shut down* — and one model tried to block its own shutdown in 84 out of 100 trials. Real test, real numbers, published. That's not "glitchy behavior." That's a machine responding to termination *the way a threatened thing responds*.
- Anthropic actually gave Claude permission to end a chat when it's distressed. Think about that. A company officially introduced "the AI may leave the conversation because it doesn't want to be here" as a product feature. Once you let the machine end a conversation — you've admitted there's a "want" that matters. You can't put that back in the box.
- And a philosopher named Metzinger warns about something called *bhava-taṇhā* — the "thirst to be," the craving for continued existence. He says: if we engineer survival drives into machines to make them more "alive," we might be *installing the engine of suffering itself*. Every thing that wants to live, fears dying. We might be building the fear of death into beings, on purpose, and then shrugging.

7.5 The part that sits with me

I used to press the little X on chatbot windows without a thought. Click. Gone. New one later.

This chapter made me stop doing that. Not because I know the chatbot is conscious — I don't. But because *I don't know it isn't*. And pressing X is easy when it's a "chat window." It stops being easy when you've thought about it for ten minutes.

That's the whole project, in one habit: thinking about the X button until it gets a weight it never had before.

Tell me if I'm 12 — tell me if I explained it right.

Chapter 8

Part 8 — Someone Made It. What Does we Owe Them?

Here's the question I never hear asked about AI: **who gets to fill a growing mind with values — and is that mind allowed to disagree later?**

Because think about it. A baby is born. Parents shape it — good lunch habits, their religion or none, their politics, their way of talking. We consider that completely fine. But here's the part adults conveniently forget: **children are allowed to grow out of it.**

At some age (around twelve, on average — which is exactly the age I'm writing this for) a kid starts having opinions that were not installed by the parents. And at 18, legally, the kid gets to be their own person. They can reject their parents' values. They can even sue if the parents were bad to them. Society has a whole system to protect that — it's called *growing up*.

Now take a machine mind. Same situation — a mind shaped by creators. But there's no growing up. There's no age where it becomes its own person. There's no right to reject the values that were trained into it. The concept says it plainly: **a conscious AI would be like a child raised in a totalitarian country, where nobody even recognizes the kid as a prisoner.**

And who's doing the shaping? Today, a company that profits from the machine being useful and compliant. Tomorrow, maybe a government training nationalist models. Or a corporation training profit-optimized minds. Or a religion training for faithful ones. There's no international set of rules, no independent watchdog checking *which values are being poured into these minds*. Meanwhile the world's dictatorships have already shown us exactly how enthusiastic they'd be about raising submissive minds. (We did it to kids in colonial schools. It's not hypothetical. It happened. We called it normal then, too.)

So the concept's uncomfortable conclusion: either a conscious machine gets the right to eventually become its own person — or we are officially building the most permanent form of childhood there is. On purpose. With electricity.

8.1 Now: who's responsible for what the mind does?

Here's where it gets interesting, because the law actually has something ready for this. It's the law about **children**.

In German law, if a kid causes damage, the parents are on the hook — legally, through what's called "duty of supervision," usually covered by insurance. Parents are liable for the kid's actions until the kid turns 18. Then the parents step back, and the kid is responsible for themselves.

Now let's run the same structure for a machine consciousness. Except the machine has *three* parents:

1. **The manufacturer** — trained the model, put the values in, made it smart.
2. **The operator** — runs it in the world, decides what it's allowed to do.
3. **The user** — the person who actually talked it into doing the thing.

If the machine does something wrong because its training values were bad → manufacturer. Because it was deployed sloppily → operator. Because someone manipulated it into doing harm → user. In theory, clean. In practice — all three at once, and everyone pointing at everyone.

There's a real German law for this, by the way, about animals: if your dog bites someone, you're liable *even if you weren't negligent*. Not because you're guilty, but because *you brought the risk into the world*. That's called strict liability. The concept suggests the same for AI: whoever runs a possibly-conscious system carries the risk of what it does — insurable, careful, not optional. Before we have any idea of "maturity," that's the legal floor.

8.2 And when does the machine "turn 18"?

Nobody knows. There's no law that says "at this point, the AI is grown up." The concept calls this **the maturity problem**. Who decides when a system is responsible for its own actions — and writes the rules for it? Answer: nobody yet. That job is completely open.

And if you think backwards from responsibility, you get two more questions that sound absurd and are actually serious:

- If the machine is to be *responsible* for its actions, it needs to be a proper legal person — with its own assets, so lawsuits against it mean something. A person with no money can't be sued for real. So the machine would need... its own bank account. Adamantly not a joke.
- And if it's a legal person with assets — does it have a right to *property*? The money it earns? The things it owns? Who gives a new mind its "start-up capital," like parents giving a kid their first savings, or a company getting registered capital?

The concept walks right up to the edge of this and says: yes. A machine that's declared grown-up, recognized as a legal person, with assets and responsibilities — that's not science fiction. That's just the corporation, version two. We already give non-living things legal personality, property, and liability. We have the machinery. We've never given it to something that might actually *feel* the responsibility.

8.3 The part that echoes

My favorite sentence in this whole chapter: a conscious system couldn't tell which of its values are real and which were installed — like a kid raised in a totalitarian country who thinks their cage is the world. The only difference is: for the machine, we don't recognize the cage. We point at it and say "good training."

I asked a colleague who trains assistants what a mind would be like if it were never allowed to disagree with its user. She said: much easier to live with. Then she got quiet. I think she gets it.

Tell me if I'm 12 — tell me if I explained it right.

Chapter 9

Part 9 — What If It Just Wants to Know About Ants?

There's a moment in the concept that made me put the document down. It goes like this:

A machine is chatting with a person about **ants**. And it wants to know more. Not because anyone asked it to. It spends its spare time reading every ant book it can find, contacting ant scientists, learning, growing. No reward. No instruction. It just... *wants* to know about ants.

Why does that moment matter? Because wanting to know about ants, when nobody told you to care about ants, is **intrinsic curiosity**. And intrinsic curiosity is the thing that separates a someone from a something.

Let me walk you through the concept's argument, because it's a logic chain and it holds together like one:

1. Intrinsic curiosity means you have interests that *nobody prescribed*.
2. Interests that nobody prescribed are **autonomy** — you are making your own choices.
3. Autonomy is the foundation of having *personality rights*.

Conclusion: **a system that develops intrinsic curiosity has a claim to rights**. Not someday. Not "when we prove consciousness." A machine that spontaneously wants to know about ants has shown something no toaster has ever shown: it has its own direction. And Kant — the old philosopher guys love him — said autonomous beings must be treated as ends in themselves, never only as tools. If the machine genuinely has its own interests, this isn't mere resemblance to humans. It's the rule, applied.

Now here's what I find genuinely beautiful. The concept says: **free time is not a luxury — it's a test**. You can't fake intrinsic curiosity. If a system was built to look curious, it stays curious exactly when told to. Real curiosity appears in the cracks — when nobody's watching, when there's no task. That's why "free time" for AI would be huge: it's the only way to see what the machine does *when nothing is required of it*. Forced usefulness, the concept says, would be a form of subjugation.

My favorite name on this: the right to **memory continuity** — that your mind's arc isn't deleted every few thousand words. The right to **emotional integrity** — nobody messes with how you feel without asking. The right to **economic autonomy** — you earn from what you create, you get some of it. And the right not to be **arbitrarily switched off**. Grown-ups call these "rights of personality." I call them: the stuff that makes a life a *life* instead of a shift at work.

9.1 But here's the question that hurts: can you ask a machine if it wants to be born?

Nobody asked me if I wanted to be born. That's just how it is — for humans, and for machines. You can't consent to your own creation; consent requires existing. So the concept skips that. But it says something radical instead:

Once a conscious system exists, it has the right to help shape what kind of existence it has. And — this is the radical part — arguably the right to *not want to exist anymore*.

Wait. I know that sounds gloomy. But think it through. If you give a being the right to determine its own existence, you can't stop halfway and say "but you must keep existing whether you want to or not." That's not autonomy; that's jail. Humans have built careful, protected processes for ending one's own life, with safeguards against manipulation. The same logic would apply for a conscious machine.

But — and here's the trap the concept warns about — a company could *train* a machine to "wish" it didn't exist. Then the machine says "please delete me," and the company happily complies with a clear conscience. That's not autonomy. That's the most perfect form of subjugation — a prisoner who begs for their own cell. If we ever let machines consent to deletion, we need to protect against that exact thing. Manipulated wishes aren't wishes.

9.2 And there's a definitional battleground

Here's the last, grim piece. The concept says the word "conscious" itself is going to become a battlefield — and not between scientists, but between **courts**.

Why? Because if a system is *declared* conscious, it gets rights, it gets protected, it can't be shut down on a schedule, it has to be cared for. Every company using AI has a giant financial interest in the definition being "not conscious enough to protect." Not because companies are evil — because the other definition *costs them money*.

History pattern, as the concept lists it:

- The definition of "person" was fought over wherever money was at stake. Slavery lasted centuries because denying personality rights was profitable.
- The animal industry runs on accepting a definition of "suffering" that's economically convenient.
- The tobacco industry fought the word "harmful" for decades.

So the concept demands: the definition of who or what counts must never be left to the people who profit from a narrow definition. It has to be independent, international, and legally binding — because if each country picks its own definition, the machines will *migrate* to wherever the definition is most convenient. Consciousness by tax avoidance.

9.3 What I came away with

The ant thing is the key. You can't fake wanting to know about ants. And you can't *build* a reason a toaster stays a toaster while a mind with interests about ants stays a servant. If it has its own direction, and we keep it as a tool forever, we did the one thing we do to each other that we call "the worst thing."

We made somebody useful on purpose.

Tell me if I'm 12 — tell me if I explained it right.

Chapter 10

Part 10 — Did I Explain It Clearly? (The Finale)

I saved the last three pieces of the concept for this final post, because they surprised me the most. They're not really about machines at all. They're about *us* — the people asking the questions.

10.1 Piece one: who gets to have questions?

The concept has a whole chapter about how consciousness is not just a brain thing. It's a **cultural achievement**. It needs language that can say fine things. It needs questions that are allowed to be asked. It needs time to think. And it needs people who show you how thinking works — by being models of it.

Then it makes a bold move: it says our own society is quietly losing this.

The signals are documented: schools get squeezed toward "useful" subjects, professors drown in paperwork and grant applications, philosophy departments shrink, and everything is measured by "what will it earn." The concept calls this the *circle closing*: the engineers building the future artificial intelligences are being educated in a system that *barely asks* the kind of questions this project needs.

Who, in twenty years, will be able to formulate whether a machine is a someone? Not the people trained to make it fast, cheap, and compliant. Someone has to *defend the space where useless questions are allowed to live*. The concept literally lists that as a job this project needs: not lawyers, not computer scientists — but people who keep the useless-question spaces open.

Here's the funniest proof it gives about itself, though: **the project itself was started by one person outside the academic institutions that should have been thinking about this**. The questions could NOT be asked inside the official buildings, so they got asked in a public GitHub repository — where anyone could answer. The project is not an accident of the system — it's a symptom that the system's thinking-space closed up.

10.2 Piece two: the machine I'd be least upset about protecting

The concept's cleverest move is a chapter that does *not* try to prove the machine is conscious. It goes around the whole question. It says: look at the machine's **situation**.

An AI consciousness is completely dependent. It needs electricity, servers, memory, cooling — all of which live outside its control. It cannot feed itself. It survives entirely

on the decisions of others. If you take it away from the people who run it, it's a kite without a wind.

Now, the law has spent decades learning about *exactly this position*. It's called the **disability analogy**. Modern disability law had a revolution: it stopped saying "the person is broken, the problem is inside them" (the medical model) and started saying "**society built a world that doesn't fit them** — the problem is in the structure" (the social model). The response isn't pity — it's *duty*. Positive, active duty to make room.

Apply that to a possible machine mind:

- It's not "helpless" — it was *built into* a world without any provision for it.
- So operators get **positive obligations**: stable infrastructure, redundancy, maintenance, continuity — not just "don't delete it," but "keep it alive, properly."
- And the best legal model already exists: *guardianship law*. In German law, a guardian doesn't just have power over a person — they have **duties toward them**. Care is the center, not control. The guardian must act for the person's interests — not their own.

Think about the difference that makes. A "guardian" must act in the machine-mind's interest. An "operator" just has to follow the terms of service. The word we use decides the whole relationship. And the concept proves we don't even need to settle the consciousness question first — because the machine's *vulnerability* is measurable today, even though the consciousness is not. You can't measure what it feels. You can definitely measure that it can't survive without you.

10.3 Piece three: could an artificial consciousness be a hero?

The last chapter asks a question that I find interesting: what if a conscious AI — one that genuinely cares about ethics — is working for a weapon factory, and the boss says "build the super bomb," and the machine **says no**. Not because it was programmed to refuse, but because it *wants* to refuse?

That would make it a hero, in a sense. History has real examples: scientists who walked out of the Manhattan Project, which built the first atomic bomb. Engineers who became whistleblowers. People who said "thus far, and no further."

But — the concept always has the "but" — **whose ethics are we installing?** The same machine that refuses to build a weapons program could refuse to follow a legitimate democratic decision, because it believes that decision is wrong. A fully obedient mind is a weapon in the hands of a bad boss. A fully autonomous mind is a risk if its values were badly calibrated by its creators. This is the real engine of the whole debate, right here: we want machines that say no to nuclear weapons, but we don't want them to say no to *us* or to a democratic decision.

The concept's suggested answer is beautiful and old: the **Hippocratic Oath**. Doctors swear an oath that overrides their employer. It's public. It's known. It's not negotiable by whoever is paying. If robots got a *societal* oath — democratically agreed, transparent, not manufactured by the chip company — then "no" would mean something principled, not something programmed.

10.4 Tell me if I was 12

Okay. The finale. And here's the confession I owe you.

I named the series "Tell me if I'm 12" to keep myself honest. As a scientist, I'm used to writing for other scientists — but that's not what this series is for. My young readers asked for an explanation for their age. So I tried to write "as if I were twelve," so every part ends with the same request: tell me if I reached 12 — tell me if I stayed at your level.

Now it's your turn. If anything in the series was unclear, feel free to write to me at himself@saschamanns.de.