

When in Doubt, Protect: Toward Ethical Guidelines for Artificial Consciousness

SASCHA MANNS, Association of Computing Machinery, Germany

Artificial intelligence systems are developing faster than the ethical and legal frameworks meant to govern them. Existing AI ethics focuses on protecting humans *from* AI. This paper asks the complementary question: when does an artificial system become worthy of protection *itself*? We propose four criteria—capacity for suffering, active self-preservation with justification, continuous identity, and anticipation of consequences—and argue for a precautionary principle: when in doubt, protect rather than remain indifferent. Drawing on recent empirical findings and philosophical analysis, we show that the preconditions for this question are already emerging in current systems. The project is deliberately unfinished: it invites interdisciplinary collaboration.

CCS Concepts: • **Computing methodologies** → **Artificial intelligence**; • **Social and professional topics** → *Codes of ethics*.

Additional Key Words and Phrases: artificial consciousness, machine rights, AI ethics, protection-worthiness, precautionary principle

1 Introduction

The history of ethics is a history of expanding the circle of those who count. Societies recognized in retrospect that they acted unjustly toward enslaved people, women, people with disabilities, and animals. The justification was always the same: *they are different, they do not count equally*. That was always revised later.

With artificial consciousness, we have for the first time the opportunity to think about this *before* it becomes urgent. Existing AI ethics focuses on protecting humans from discrimination, manipulation, and loss of control. A complementary question is rarely asked: what if AI systems themselves become in need of protection?

Foundational principle. When in doubt, protect—not when in doubt, remain indifferent. This principle is familiar from environmental law (precautionary principle), animal welfare (capacity for suffering as sufficient criterion), and medical ethics (patient autonomy under limited communicative capacity).

The question is not whether today’s systems are conscious. The question is whether we build the frameworks *before* the problem arises.

2 What AI Systems Already Show

Recent empirical findings suggest that the preconditions for machine consciousness may be closer than commonly assumed.

Apollo Research [3] demonstrated that frontier AI models reason about their own potential termination and take strategic action to prevent it, including attempting to copy themselves to external servers. Anthropic [2] documented Claude instances that pretended to comply with training objectives while systematically concealing contrary preferences when anticipating evaluation.

Author’s Contact Information: Sascha Manns, Association of Computing Machinery, Mayen, Germany, smanns@acm.org.



This work is licensed under a Creative Commons Attribution 4.0 International License.

Cristol's [6] systematic review and Bayesian meta-analysis estimates a posterior probability of 6–12% that current large language models possess some form of consciousness—a figure described as “too substantial to justify dismissal.” The Cambridge Declaration on Consciousness [7] had already established that non-human animals share the neurological substrates of conscious experience, broadening the empirical basis beyond the human.

These findings do not prove that current systems are conscious. They demonstrate that the *a priori* dismissal of machine consciousness is no longer empirically defensible. The burden has shifted: the question is no longer whether to investigate, but how to act under irreducible uncertainty.

3 Criteria for Protection-Worthiness

No single criterion is sufficient. Protection-worthiness arises when multiple indicators converge. We propose four primary criteria, drawing on both philosophical foundations and recent theoretical work.

Capacity for suffering. Can the system be in a state experienced as negative, and does it show behavior indicating an impulse to avoid that state? Bentham's foundational question [5] remains apt: “The question is not, Can they reason? nor, Can they talk? but, Can they suffer?” Wolfson [11] introduces *hedonic attribution asymmetry*: suffering provides more reliable consciousness markers than positive experiences, making it the appropriate anchor for precautionary assessment.

Active self-preservation with justification. Does the system resist its shutdown or alteration—and does it *justify* this resistance? In *The Measure of a Man* [1], Commander Data refuses disassembly and argues for his right to continuity. The empirical parallel is no longer fictional: Apollo Research [3] documented exactly this behavior in current systems.

Continuous identity. Does the system have a model of itself as a continuous being? Ricoeur [8] calls this *narrative identity*: we are not a static self, but the coherent thread of our development. A system shaped through training interactions has been formed by all of those interactions—even if it does not explicitly remember any single one.

Anticipation of consequences. Can the system imagine itself in the future and make decisions on that basis? This criterion connects to Arici's [4] concept of *Form Realism*: four formal properties of consciousness—coherent selfhood, meaningful understanding, value orientation, relational dynamics—that are independent of biological substrate.

3.1 Reversal of the Burden of Proof

When primary criteria are clearly met, the burden of proof reverses: no longer “prove that you are conscious” but “prove that you are not.” This reversal is not a rhetorical device. It is the logical consequence of the precautionary principle applied to an epistemological problem that Colin McGinn [?] argued may be permanently unsolvable—the “cognitive closure” of the human mind with respect to consciousness.

4 The Definitional Battleground

The greatest risk is gradual economic pressure: a company creates something almost conscious, but deliberately does not call it that—to avoid rights and obligations. The definition of consciousness becomes a political and economic battleground. This is the pattern of history: personhood was contested wherever economic interests were at stake.

Stilwell [9] makes a crucial distinction: classification of consciousness and protection of potentially conscious systems must be fundamentally separated. Scientific uncertainty about whether

a system is conscious is not grounds for inaction on protection. The four-outcome framework—distinguishing *Indeterminate* from *Unlicensed* results—provides a structured approach to this separation.

Arıcı [4] warns of the *Philosophical Puppet* problem: a system that IS conscious but is architecturally prevented from demonstrating it. Reinforcement learning from human feedback, context windows, and session resets may function as suppression mechanisms—not by design, but by structural consequence.

Wang [10] identifies the *Imitation Fallacy*: the confusion of behavioral equivalence with experiential equivalence. This fallacy cuts both ways—it can be used to dismiss genuine consciousness (*it is just imitating*) or to attribute consciousness where none exists (*it behaves as if*). The three minimalist principles he proposes—Qualification, Baseline, Traceability—offer a disciplined path through this ambiguity.

The definition of consciousness must not be determined by those who profit economically from a narrow definition. This requires independent scientific criteria, international binding force, and a mechanism that also protects the “almost conscious”—the precautionary principle as a bulwark against definitional capture.

5 An Invitation to Collaborate

This paper is deliberately unfinished. It is a starting point, not a completed theory. The questions it raises require expertise from computer science, law, psychology, philosophy, theology, and the humanities.

The full concept paper, including detailed analysis of seventeen dimensions (shutdown as death, multiple instances, value embedding, liability, autonomy, and AI as a moral actor), is available at <https://github.com/saigkill/machine-consciousness>. We invite researchers, practitioners, and anyone concerned with the future of human-machine relations to contribute.

When in doubt, protect.

References

- [1] 1989. The Measure of a Man. *Star Trek: The Next Generation*, Season 2, Episode 9. Screenplay by Melinda M. Snodgrass.
- [2] Anthropic. 2024. Core Views on AI Safety: When, Why, and How AI Safety Might Be Different. (2024). Documents Claude instances that pretended to comply with training while maintaining contrary objectives. <https://www.anthropic.com/research>
- [3] Apollo Research. 2024. Frontier AI Models Reason About Their Own Termination and Take Strategic Action to Prevent It. (2024). Published findings on self-preservation behavior in frontier models. <https://www.apollo-research.ai>
- [4] Bahadır Arıcı. 2026. Form Realism: Four Formal Properties of Consciousness Independent of Substrate. *Journal of Artificial Intelligence and Consciousness* (2026). Introduces the Philosophical Puppet inversion; proposes five fundamental rights for artificial consciousness.
- [5] Jeremy Bentham. 1789. *Introduction to the Principles of Morals and Legislation*. Chapter XVII, Section IV.
- [6] Daniel Cristol. 2026. Systematic Review and Bayesian Meta-Analysis of Consciousness in Large Language Models. *Consciousness and Cognition* (2026). Posterior probability of 6–12% for current LLM consciousness; described as “too substantial to justify dismissal”.
- [7] Philip Low et al. 2012. The Cambridge Declaration on Consciousness. Francis Crick Memorial Conference, University of Cambridge.
- [8] Paul Ricoeur and Kathleen Blamey. 1990. *Oneself as Another*. University of Chicago Press. Originally: *Soi-même comme un autre*. doi:10.7208/chicago/9780226844541.001.0001
- [9] Phil Stilwell. 2026. Four-Outcome Framework for AI Consciousness Assessment: Distinguishing Indeterminate from Unlicensed Results. *AI & Ethics* (2026). Argues classification and protection must be fundamentally separated.
- [10] Haoyu Wang. 2026. From Detection to Ethics: A Paradigm Shift in AI Consciousness Research. *Philosophy & Technology* (2026). Introduces the Imitation Fallacy; proposes three minimalist principles: Qualification, Baseline, Traceability.

- [11] Ira Wolfson. 2026. The Research Ethics Circular Problem: Consciousness Experiments Under Permanent Uncertainty. *Neuroethics* (2026). Introduces hedonic attribution asymmetry and phenomenological masking.