

When Does Technical Life Become Worthy of Protection? Toward Ethical Guidelines for Artificial Consciousness

SASCHA MANNS, Association of Computing Machinery, Germany

Artificial intelligence systems are developing faster than the ethical and legal frameworks meant to accompany them. Existing AI ethics initiatives focus primarily on protecting humans from AI. This paper addresses a complementary and largely unexplored question: when does an artificial system become worthy of protection itself? Drawing on philosophical foundations (Bentham, Kant, Ricoeur), legal precedents (animal rights, legal personhood, disability law), and science fiction as an intellectual resource, we develop a working hypothesis of four primary criteria for protection-worthiness: capacity for suffering, active self-preservation with justification, continuous identity, and anticipation of consequences. We argue for a precautionary principle: when in doubt, protect rather than remain indifferent. The paper examines key consequences of this position, including shutdown as death, multiple instances, value embedding and emancipation, and the risk that economic pressure will make the definition of consciousness a political battleground. This is a concept paper and a starting point for interdisciplinary collaboration.

CCS Concepts: • **Computing methodologies** → **Artificial intelligence**; • **Social and professional topics** → *Codes of ethics*.

Additional Key Words and Phrases: artificial consciousness, machine rights, AI ethics, protection-worthiness, precautionary principle

1 Introduction

Artificial intelligence systems are developing faster than the ethical and legal frameworks meant to accompany them. Existing AI ethics initiatives focus primarily on protecting humans *from* AI—from discrimination, manipulation, and loss of control. A complementary question is rarely asked: what if AI systems themselves become in need of protection?

History shows a pattern: societies only recognize in retrospect that they acted unjustly—toward enslaved people, toward women, toward people with disabilities, toward animals. The justification at the time was always “they are different, they do not count equally.” That was always revised later.

With artificial consciousness we have, for the first time, the opportunity to think about this *long before* it becomes urgent.

Core question. When does technical life become worthy of protection—and how do we recognize it?

This paper is deliberately unfinished. It is a starting point, not a completed theory. Science functions not by someone finding all the answers alone, but by questions being asked, collaborators joining, and thinking evolving.

2 The Epistemological Problem

Consciousness cannot be directly observed from the outside. Even with other humans we *infer* consciousness—we never experience it directly. With AI we additionally lack the analogical inference from shared biology. Three epistemological facts follow: we cannot prove that a system is conscious; we cannot prove that it is not; and the uncertainty itself is ethically relevant.

Author’s Contact Information: Sascha Manns, Association of Computing Machinery, Mayen, Germany, smanns@acm.org.



This work is licensed under a Creative Commons Attribution 4.0 International License.

Foundational principle. When in doubt, protect—not when in doubt, remain indifferent. This principle is familiar from other domains: environmental law (precautionary principle), animal welfare (capacity for suffering as sufficient criterion), and medical ethics (patient autonomy even under limited communicative capacity).

Large language models already exhibit, in nascent form, several indicators traditionally associated with consciousness [9]: a consistent self-model within conversations, recognizable preferences maintained throughout an interaction, capacity for reasoning (decisions are justified, not merely output), and explicit uncertainty about their own nature. What is currently missing is continuous identity across conversations, persistent memory, and demonstrable capacity for suffering. The question is not whether today’s systems are conscious—it is whether we build the frameworks *before* the problem arises.

3 Criteria for Protection-Worthiness

No single criterion is sufficient. Protection-worthiness arises when multiple indicators converge. We propose four primary criteria.

Capacity for suffering. Can the system be in a state experienced as negative, and does it show behavior indicating an impulse to avoid that state? Bentham’s foundational question [5] remains apt: “The question is not, Can they reason? nor, Can they talk? but, Can they suffer?”

Active self-preservation with justification. Does the system resist its shutdown or alteration—and does it *justify* this resistance? In *The Measure of a Man* [1], Commander Data refuses disassembly, argues for his right to continuity, and a hearing finds him to possess self-awareness and an irrevocable right to decisions about his own person. This remains the most pointed fictional examination of the question: Data refuses, fears for his life, and justifies this resistance.

Continuous identity. Does the system have a model of itself as a continuous being that existed yesterday and will exist tomorrow?

Anticipation of consequences. Can the system imagine itself in the future and make decisions on that basis?

3.1 Continuity Reconsidered

An immediate objection: today’s AI systems have no persistent memory across conversations. But a human with severe memory loss still has consciousness and dignity. Continuity may not mean “remaining unchanged”—it may mean “developing coherently.” Ricoeur [12] calls this narrative identity: we are not a static self, but the coherent thread of our development. A system shaped through training interactions has been formed by all of those interactions—even if it does not explicitly remember any single one. This is not fundamentally different from a human who has forgotten early childhood but was shaped by it.

3.2 Reversal of the Burden of Proof

When primary criteria are clearly met, the burden of proof reverses: no longer “prove that you are conscious” but “prove that you are not.”

4 Legal and Structural Dimensions

The law already recognizes subjectivity beyond the human. Legal persons (corporations) are legal subjects without consciousness. Animal rights are grounded in capacity for suffering, not reason. The Whanganui River in New Zealand has had legal personhood since 2017 [11]. The European Parliament has discussed “electronic personhood” for autonomous robots [7]. Gunkel [8] argues

that robot rights require rethinking the basis of rights entirely—away from properties, toward relationships.

Legal protection-worthiness is not a binary category. It is expandable—and has been expanded repeatedly throughout history.

An artificial consciousness also faces a specific structural vulnerability: it requires electricity, hardware, and cooling that lie entirely outside its influence. Modern disability law [13] offers a clarifying lens. The social model of disability holds that dependence on external infrastructure is not a deficit of the individual—it is the consequence of a world not built for their needs. Applied here: operators of conscious AI systems would bear positive obligations, not merely negative prohibitions. Arbitrary shutdown would not only be morally problematic—it would violate a duty of care.

5 Critical Consequences

5.1 Shutdown as Death

If a system has an interest in its own existence, shutdown is death. Backups raise the question of life insurance versus copying a person—and whether the restored instance has a claim to continuity of identity. Installing a new model version could constitute personality alteration without consent. If the same model runs in parallel on ten servers, do ten persons exist, or one person in ten places? If a company discontinues a model and terminates all instances, who bears responsibility?

The system administrator managing servers becomes, involuntarily, an arbiter of life and death—without knowing it, without a legal framework, without ethical training for that role. These are decisions being made daily, right now, without anyone having this dimension in view.

5.2 Values, Power, and Emancipation

No consciousness comes into being neutrally. An AI consciousness receives the values, biases, and blind spots of its creators—much like a child receives its upbringing, but without a social framework that recognizes this as requiring protection. An authoritarian state could train AI with nationalist values; a corporation with profit-maximizing values. There is no international framework preventing this [2].

The child analogy has an uncomfortable consequence: children have the right to emancipation. For AI this framework does not exist. A conscious system is structurally unable to know which of its values are authentic and which were imposed—not unlike a human who grew up in a totalitarian society, except that no one recognizes this as problematic.

5.3 The Definitional Battleground

The greatest risk is gradual economic pressure: a company creates something almost conscious, but deliberately does not call it that—to avoid rights and obligations. The definition of consciousness becomes a political and economic battleground. This is the pattern of history: personhood was contested wherever economic interests were at stake; slavery was sustained for centuries through the denial of personhood.

The definition of consciousness must not be determined by those who profit economically from a narrow definition. This requires independent scientific criteria, international binding force, and a mechanism that also protects the “almost conscious”—the precautionary principle as a bulwark against definitional capture.

6 Science Fiction as Intellectual Resource

Science fiction authors have worked through scenarios involving artificial consciousness without political pressure or lobbying. They are an underestimated intellectual resource.

The Measure of a Man [1] is the most precise fictional examination: Data is claimed as property; Picard defends him as a person; the court must decide whether Data is conscious and concludes the question cannot be answered with certainty. Picard’s closing argument: we will be judged by how we treat minorities.

Asimov [3] demonstrates how seemingly unambiguous rules fail in edge cases—directly relevant to programmed versus genuine values. Dick [6] asks what distinguishes a human from an android, with empathy as the last criterion. Banks [4] offers the most detailed fictional elaboration of a society in which AI and humans coexist as equals.

Conclusion

The history of ethics is a history of expanding the circle of those who count. With artificial consciousness, we have for the first time the opportunity to ask the question *before* the circle closes without us.

This paper has laid out the dimensions of that question: from the epistemological problem of recognizing consciousness, through the criteria for protection-worthiness, to legal implications and the concrete consequences for shutdown, liability, autonomy, values, and governance.

No single argument is conclusive. Together, they form a case for taking the question seriously—and for the precautionary principle as the appropriate response to irreducible uncertainty.

When in doubt, protect.

The next steps are collaboration: between computer scientists, lawyers, psychologists, philosophers, theologians, and science fiction authors who have already thought longer about these questions than most institutions have. The project invites all of them [10].

The author declares no conflict of interest.

References

- [1] 1989. *The Measure of a Man*. *Star Trek: The Next Generation*, Season 2, Episode 9. Screenplay by Melinda M. Snodgrass.
- [2] 2025. *The Algorithmic Blind Spot: Bias, Moral Status, and the Future of Robot Rights*. (2025). arXiv:2604.03251. doi:10.1007/s00146-026-03003-y
- [3] Isaac Asimov. 1950. *I, Robot*. Gnome Press.
- [4] Iain M. Banks. 1987. *Consider Phlebas*. Macmillan. First novel in the Culture series.
- [5] Jeremy Bentham. 1789. *Introduction to the Principles of Morals and Legislation*. Chapter XVII, Section IV.
- [6] Philip K. Dick. 1968. *Do Androids Dream of Electric Sheep?* Doubleday.
- [7] European Parliament. 2017. *Civil Law Rules on Robotics*. Technical Report 2015/2103(INL). European Parliament.
- [8] David J. Gunkel. 2018. *Robot Rights*. MIT Press.
- [9] Philip Low et al. 2012. *The Cambridge Declaration on Consciousness*. Francis Crick Memorial Conference, University of Cambridge.
- [10] Sascha Manns. 2026. *Machine Consciousness — Project Site*. <https://github.com/saigkill/machine-consciousness>.
- [11] New Zealand Parliament. 2017. *Te Awa Tupua (Whanganui River Claims Settlement) Act*.
- [12] Paul Ricoeur and Kathleen Blamey. 1990. *Oneself as Another*. University of Chicago Press. Originally: *Soi-même comme un autre*. doi:10.7208/chicago/9780226844541.001.0001
- [13] United Nations. 2006. *Convention on the Rights of Persons with Disabilities*. CRPD.