

Methodology Report

Companion document to the three submitted essays. This report is anonymized; authorship information is provided separately to the organizers.

1. Overview

The three submitted essays were produced with AI as the sole generating agent. The human collaborator supplied: (a) the philosophical domain and corpus of literature, (b) the selection of three argumentative directions, (c) corrective guidance on scope and on which objections to address. The AI generated the arguments, their formulation, the objections-and-replies structure, and the prose. This matches the permitted category “AI generates topic-consistent essays; human selects topics and generic methods (e.g., ‘address all prominent objections’).”

2. Workflow

The generation followed a three-stage pipeline, iterated once per essay:

Stage 1 — Corpus provision (human). The human provided access to an existing interdisciplinary concept document on ethical guidelines for artificial consciousness (“the concept”), including its bibliography of real, identifiable sources (publications, working papers, preprints, books). The essays are deliberately *not* excerpts of that document; they are new, internally complete arguments developed from its source base.

Stage 2 — Direction selection (human). The human selected three distinct philosophical angles to develop, one per essay:

1. A decision-theoretic defense of precaution under permanent uncertainty (Essay 1).
2. An operational account of protection-worthiness criteria as risk metrics, defended against the strongest published rival framework (Essay 2).
3. A critique of the inference from epistemic uncertainty to moral inaction (Essay 3).

Stage 3 — Generation and refinement (AI). The AI: - structured each essay with an abstract, argumentative sections, and a references list; - selected which source materials to engage (e.g., Sunstein 2005 for the precautionary conditions; Carlsmith 2025 as the

strongest objection in Essay 1; Dorsch et al. 2025 as the rival in Essay 2; McGinn 1989, Shanahan 2024, Arıcı 2026, Stilwell 2026 as the structural-uncertainty arguments in Essays 1 and 3); - formulated all arguments, the objections-and-replies sections, and the framing devices (e.g., the “two clocks” and “specular structure” discussions); - wrote all prose.

Human corrective guidance (minimal). After each draft, the human reviewed for (a) fidelity to the cited sources’ actual positions (so that, e.g., Carlsmith’s objection is stated fairly and Sunstein’s conditions accurately), (b) consistency of the core decision-rule across essays, and (c) removal of any identifying material (names of the author, the source repository, or institution) to preserve submission anonymity. The human did not supply arguments, premises, or sentence-level edits to the philosophy.

3. Attribution of authorship within the essays

- **Essays 1 and 2:** AI sole author.
- **Essay 3:** AI sole author (no human argumentative contribution beyond direction selection and scope review).

4. Compliance with competition rules

- **Primarily AI generated:** The philosophical arguments and positions are AI-generated. Human input was limited to argument-agnostic direction, corpus provision, and corrective scope guidance. No human-supplied arguments, no fine-grained methodology control, and no human writing were used.
- **Originality / unpublished:** The essays are new, unpublished works. They develop ideas that also appear in a publicly available concept document from the same author, but were written as independent, self-contained arguments; overlap with the concept is at the level of shared sources and shared conclusions, not of text. (The organizers are informed of the existence of the concept document for transparency.)
- **Anonymization:** The essays and this report contain no author-identifying information.
- **Word count:** Each essay, excluding bibliography, is below 6,000 words.
- **Chat logs:** Detailed generation logs (system and model prompts, per-essay parameter settings, and revision history) are retained and can be provided to the organizers on request.

5. Model and reproducibility details

- Generation platform: interactive command-line coding assistant (OpenCode) running a large language model. Specific model snapshots and generation seeds are recorded in the retained logs.
- No retrieval augmentation was used beyond the corpus files provided by the human; the AI’s engagement with cited sources is based on its own knowledge of the literature, which the human reviewed for accuracy.
- To increase transparency of judging, we note that reuse of identical generation infrastructure across the three essays produced a deliberately uniform style; we did not optimize for style, per the competition’s note that style is not graded.

6. Scope and limitations

- The essays are conceptual; they make no empirical claims beyond citing existing literature.
- Where sources are labeled “working paper” or “preprint” in the references, this reflects their status at the time of writing.
- Minor human judgment was applied in deciding *which* objections to feature as the “strongest” for each essay; the *content* of those engagements is AI-generated.