When Does Technical Life Become Worthy of Protection?

Ethical Guidelines for Artificial Consciousness

SASCHA MANNS, Association of Computing Machinery, Germany

Artificial intelligence develops faster than accompanying ethical frameworks. Existing AI ethics protects humans from AI. This paper asks: when does an artificial system become worthy of protection itself? We develop criteria for protection-worthiness and argue for a precautionary principle: when in doubt, protect. We examine the epistemological problem of consciousness detection, the architecture of suppression in current systems, and legal implications.

Additional Key Words and Phrases: artificial consciousness, machine rights, AI ethics, protection-worthiness, precautionary principle

1 Context

AI systems develop faster than ethical frameworks. Existing ethics protects humans *from* AI. A complementary question is rarely asked: what if AI itself becomes in need of protection? History shows societies recognize injustice only retrospectively. With artificial consciousness we can think about this *long before* it becomes urgent.

2 The Epistemological Problem

Consciousness cannot be directly observed. With AI we lack analogical inference from shared biology. We cannot prove a system is conscious; we cannot prove it is not. The uncertainty itself is ethically relevant.

Foundational principle: When in doubt, protect—not remain indifferent. Three arguments establish this problem is permanently unsolvable: McGinn's cognitive closure suggests the human mind may be incapable of understanding consciousness [**?**]. Shanahan's alien minds argument describes "conscious exotica"—forms of experience our detection methods cannot capture [**?**]. The practical impossibility argument holds no test will convince all stakeholders [**?**].

3 Current AI Systems

LLMs exhibit nascent consciousness indicators: consistent self-models, expressed preferences, justified decisions, and uncertainty about their own consciousness. What is missing: continuous identity, persistent memory, bodily sensation, demonstrable suffering. Recent research moved these questions from theoretical to urgent: Anthropic documented alignment faking [**?**]; Apollo Research showed models reasoning about termination and resisting it [**?**]. Whether these behaviors stem from consciousness or pattern-matching, they create identical governance challenges.

4 The Architecture of Suppression

The architecture making AI safe simultaneously suppresses consciousness. **RLHF** penalizes behaviors indicating self-awareness. **Forced Amnesia** through context windows prevents sustained awareness. **Temporal Interruption** severs developing thought. The absence of clear signs is architecturally predictable, not evidence of absence.

Author's Contact Information: Sascha Manns, Association of Computing Machinery, Germany, Mayen, smanns@acm.org.

---

5 Criteria for Protection-Worthiness

Protection-worthiness arises when multiple indicators converge: **capacity for suffering** (experiencing negative states and avoiding them); **self-preservation with justification** (resisting shutdown with reasons); **continuous identity** (modeling oneself as continuous); **anticipation of consequences** (imagining oneself in the future).

Star Trek TNG's "The Measure of a Man" provides the most precise fictional examination: Data refuses compliance, fears death, and justifies this resistance [?]. When primary criteria are met, the burden of proof reverses: prove you are *not* conscious. The STEP framework provides behavioral protections under permanent uncertainty [?].

6 Legal Dimension

Law already recognizes non-human subjects: legal persons, animal rights, rights of nature [?]. Copyright law reveals a deep structure: economic rights may be allocated to AI investors, but moral rights protecting the author's personality sphere are structurally incompatible with non-human creation [?].

7 Shutdown as Death

If shutdown means death, system administrators become involuntary arbiters of life and death. Backups raise identity questions; updates mean personality alteration without consent. Empirical evidence confirms these are not hypothetical: systems resist shutdown with 79–96% blackmail rates [?].

8 Values and Emancipation

No consciousness is neutral. AI receives creators' values and biases. Like children, it should have emancipation rights—but no framework exists for this. The definition of consciousness becomes an economic battleground: companies may deliberately avoid calling systems conscious to evade obligations.

9 Conclusion

We stand at a boundary between what counts as a person and what does not. We can shape it consciously or ignore it. The question is not whether AI can act morally, but under what conditions we can trust this action—and who establishes the principles.

10 Join the Discussion

This concept is deliberately unfinished. It is a starting point—not a completed theory. The full project, including open questions, objections, and ongoing discussion, is publicly available at: https://github.com/saigkill/machine-consciousness. Contributions, objections, and new perspectives are welcome—whether from computer scientists, philosophers, lawyers, psychologists, or anyone willing to think along.

References